Condor: A Neural Connection Network for Enhanced Attention

Youngseong Kim

August 2025

Abstract

The attention mechanism in traditional neural networks relies on pairwise interactions between tokens, limiting its ability to capture complex, multi-token relationships. This study introduces Condor, a novel architecture that extends the attention mechanism through a neural connection network based on the KY Transform theory. Our approach replaces static attention patterns with learnable connection functions that dynamically model relationships within a local window. The Condor architecture achieves linear computational complexity of O(LWH) while maintaining the expressive power to capture sophisticated sequence patterns. Experimental results on WikiText-2 demonstrate improved perplexity and faster convergence compared to the standard Transformer, confirming that each attention head learns unique connection patterns specializing in different aspects of sequence modeling. Code is available at: `https://github.com/Kim-Ai-gpu/Condor`

1 Introduction

The attention mechanism has become the cornerstone of modern neural network architectures, especially in sequence modeling tasks. Since the introduction of the Transformer [14], self-attention has proven highly effective at capturing long-range dependencies and contextual relationships in sequences. However, traditional attention mechanisms are fundamentally constrained by the pairwise interaction paradigm, which calculates attention weights based solely on the similarity between pairs of tokens.

This pairwise constraint manifests in several ways. First, standard attention cannot directly model complex relationships involving multiple tokens simultaneously. Second, the attention pattern is determined by a fixed mathematical operation (typically a dot product followed by a softmax), which limits its adaptability to diverse sequence structures. Third, its quadratic complexity with respect to sequence length, $O(L^2)$, makes it computationally impractical for very long sequences.

Recent studies have proposed various approaches to overcome these limitations. Linear attention attempted to reduce computational complexity but at the cost of expressive power, while sparse attention patterns were limited to specific, predefined patterns. Local attention mechanisms improved computational efficiency but still operated within the fundamental framework of pairwise interactions.

In this paper, we present a completely new approach to address these fundamental limitations. Based on a mathematical theory called the KY Transform, we introduce

1

the concept of implementing an arbitrary function that connects two points as a learnable neural network. This allows us to generalize the traditional calculus concept of a "straight-line connection" to a "learnable connection function" and apply it to the attention mechanism.

The main contributions of this study are as follows:

- Mathematical definition of the KY Transform theory and its application to attention.

- A methodology for dynamically modeling token relationships through learnable connection functions.

- The design of the Condor architecture, achieving O(LWH) linear complexity.

- Multi-scale representation where each attention head learns a different connection pattern.

- Theoretical proof of the expressiveness and convergence of Neural KY-Attention.

- Experimental validation of performance improvements on the WikiText-2 dataset.

2 Related Work

2.1 Evolution of Attention Mechanisms

The self-attention in the Transformer architecture brought about breakthrough performance improvements in various natural language processing tasks, including machine translation. However, its computational complexity, which is quadratic in sequence length, has been a significant bottleneck for processing long sequences.

Various approaches have been proposed to address this. Linformer [15] reduced complexity by projecting the attention matrix into a lower dimension, and Performer [3] achieved linear complexity using kernel methods. However, these are all approximation methods and do not fully preserve the expressive power of the original attention mechanism.

2.2 Local Attention Approaches

Local attention is a method that reduces computational complexity by having each token attend only to tokens within a limited window. Longformer [2] and BigBird [16] successfully applied this approach. However, the basic framework of pairwise interaction was still maintained.

2.3 Learnable Positional Representations

The learnable extension of positional encoding has also been a significant research direction. RoPE [13] and ALiBi [9] proposed methods to more effectively model relative positional information. Our approach integrates these positional representations into the learning of connection functions.

3 Kim-Youngseong Transform Theory

3.1 Core Concept and Motivation

The traditional concept of differentiation analyzes the local change of a function through linear approximation. However, many physical phenomena and mathematical structures follow more complex, non-linear patterns. The KY Transform proposes a new concept of differentiation that extends the **"straight-line connection"** to an **"arbitrary function connection"** to overcome these limitations.

Definition 3.1 (Basic Idea of the KY Transform). *The method of connecting two points* $(x, f(x))$ *and* $(x + h, f(x + h))$:

- ***Traditional Derivative****: Connects with a straight line* $\rightarrow$ *Extracts only slope information.*

- ***KY Transform****: Connects with an arbitrary parameterized function* $\rightarrow$ *Extracts multi-dimensional information.*

3.2 Mathematical Definition and Uniqueness Condition

Definition 3.2 (Connection Function and Parameter Space). *A connection function* $g(t; \boldsymbol{\theta}) : \mathbb{R} \times \Theta \rightarrow \mathbb{R}$ *with a parameter vector* $\boldsymbol{\theta} = (\theta_1, \theta_2, \ldots, \theta_n) \in \Theta \subseteq \mathbb{R}^n$ *satisfies the following:*

1. *$g(t; \boldsymbol{\theta})$ is continuously differentiable with respect to $\boldsymbol{\theta}$.*

2. *For each t, $g(t; \cdot) : \Theta \rightarrow \mathbb{R}$ is an injective function.*

Definition 3.3 (Constraint System). *For a function* $f : \mathbb{R} \rightarrow \mathbb{R}$, *a point* x_0, *and an increment* $h \neq 0$, *the parameter vector* $\boldsymbol{\theta}(h)$ *satisfies the following basic constraints:*

$$\mathcal{C}_1(h) : \quad g(x_0; \boldsymbol{\theta}(h)) = f(x_0) \tag{1}$$
$$\mathcal{C}_2(h) : \quad g(x_0 + h; \boldsymbol{\theta}(h)) = f(x_0 + h) \tag{2}$$

Definition 3.4 (Regularization Conditions). *To ensure the uniqueness of the parameters, the following regularization conditions* $\mathcal{R}$ *are added depending on the type of connection function:*

KY-Polynomial (n-th degree):

$$\mathcal{R}_P^n : \quad \|\boldsymbol{\theta}(h)\|_2 = \min_{\tilde{\boldsymbol{\theta}} \in \mathcal{S}(h)} \|\tilde{\boldsymbol{\theta}}\|_2 \tag{3}$$

where $\mathcal{S}(h) = \{\boldsymbol{\theta} : \mathcal{C}_1(h) \wedge \mathcal{C}_2(h)\}$ *is the solution set satisfying the constraints.*

KY-Exponential:

$$\mathcal{R}_E : \quad |B| = \min_{\tilde{B} \in \mathcal{S}_B(h)} |\tilde{B}|, \quad |A| \leq M \tag{4}$$

where $M > 0$ *is an amplitude upper bound, and* $\mathcal{S}_B(h)$ *is the valid solution set for the B parameter.*

KY-Harmonic:

$$\mathcal{R}_H : \quad |\omega| = \min_{\tilde{\omega} \in \mathcal{S}_\omega(h)} |\tilde{\omega}|, \quad \phi \in [0, 2\pi) \tag{5}$$

Theorem 3.5 (Uniqueness of Parameters). *When the connection function $g(t; \boldsymbol{\theta})$ satisfies the above conditions and a regularization condition $\mathcal{R}$ is given, the parameter vector $\boldsymbol{\theta}(h)$ is uniquely determined for each $h \neq 0$.*

Proof Sketch. The constraints $\mathcal{C}_1(h), \mathcal{C}_2(h)$ provide 2 equality conditions in an n-dimensional parameter space. If $n > 2$, the solution set $\mathcal{S}(h)$ generally forms an $(n-2)$-dimensional manifold. The regularization condition $\mathcal{R}$ acts as a selection function that chooses a unique optimal solution from this manifold. The uniqueness of the solution is guaranteed by the continuity of the connection function and the independence of the constraints. $\quad\square$

Definition 3.6 (Limit Convergence of the KY Transform). *Conditions for the parameter vector $\boldsymbol{\theta}(h)$ to converge as $h \to 0$:*

1. ***Continuity Condition***: $\lim_{h \to 0} \boldsymbol{\theta}(h)$ *exists.*

2. ***Consistency Condition***: $\lim_{h \to 0} g(x_0; \boldsymbol{\theta}(h)) = f(x_0)$

3. ***Differentiability Condition***: *f is differentiable to a sufficiently high order at x_0.*

Definition 3.7 (KY Transform). *If a function $f : \mathbb{R} \to \mathbb{R}$ is KY-differentiable at a point x_0, the KY Transform at point x_0 is defined as:*

$$\nabla_{KY}[f(x); g](x_0) := \lim_{h \to 0} \boldsymbol{\theta}(h) \tag{6}$$

where $\boldsymbol{\theta}(h)$ is the unique solution that satisfies the constraints $\mathcal{C}_1(h), \mathcal{C}_2(h)$ and the regularization condition $\mathcal{R}$.

3.3 Notation System

We use the following dedicated symbols according to the type of connection function:

Definition 3.8 (Notation by Connection Function).

$$\nabla_L[f](x_0) := \textit{Linear KY-Transform (same as conventional derivative)} \tag{7}$$
$$\nabla_P^n[f](x_0) := \textit{n-th degree Polynomial KY-Transform} \tag{8}$$
$$\nabla_E[f](x_0) := \textit{Exponential KY-Transform} \tag{9}$$
$$\nabla_T[f](x_0) := \textit{Trigonometric KY-Transform} \tag{10}$$
$$\nabla_H[f](x_0) := \textit{Harmonic KY-Transform} \tag{11}$$

Access to the parameter results is notated as follows:

$$[\nabla_P^2[f](x_0)]_k = k\text{-th degree coefficient}, \quad k = 0, 1, 2 \tag{12}$$

Aliases based on meaning clarify the physical interpretation:

$$\text{curvature}(f, x_0) := [\nabla_P^2[f](x_0)]_2 \tag{13}$$
$$\text{slope}(f, x_0) := [\nabla_P^2[f](x_0)]_1 \tag{14}$$
$$\text{value}(f, x_0) := [\nabla_P^2[f](x_0)]_0 \tag{15}$$

3.4 Key Connection Functions and Their Properties

3.4.1 KY-Linear Transform

Definition 3.9 (KY-Linear). *Connection function:* $g(t) = a(t - x_0) + b$
From the constraints:

$$b = f(x_0) \tag{16}$$

$$ah + b = f(x_0 + h) \tag{17}$$

$$\Rightarrow a = \frac{f(x_0 + h) - f(x_0)}{h} \tag{18}$$

Result: $\nabla_L[f](x_0) = \lim_{h \to 0} \left(\frac{f(x_0+h)-f(x_0)}{h}, f(x_0) \right) = (f'(x_0), f(x_0))$

3.4.2 KY-Polynomial Transform

Definition 3.10 (KY-Polynomial). *Connection function:* $g(t) = a_n(t-x_0)^n + \cdots + a_1(t - x_0) + a_0$
The parameters are uniquely determined by the constraints and regularization conditions:

$$\nabla_P^n[f](x_0) = \lim_{h \to 0} (a_n(h), a_{n-1}(h), \ldots, a_1(h), a_0(h)) \tag{19}$$

In the special case of degree 2, the relationship with the Taylor expansion is:

$$\nabla_P^2[f](x_0) = \left(\frac{f''(x_0)}{2!}, f'(x_0), f(x_0) \right) \tag{20}$$

3.4.3 KY-Exponential Transform

Definition 3.11 (KY-Exponential). *Connection function:* $g(t) = A \cdot \exp(B(t - x_0)) + C$
Constraints:

$$A + C = f(x_0) \tag{21}$$

$$Ae^{Bh} + C = f(x_0 + h) \tag{22}$$

The parameters are uniquely determined by the regularization condition $\mathcal{R}_E$:

$$\nabla_E[f](x_0) = \lim_{h \to 0} (A(h), B(h), C(h)) \tag{23}$$

Physical meaning:

- *A: amplitude*

- *B: growth rate*

- *C: baseline*

3.5 Convergence Analysis and Existence Conditions

Theorem 3.12 (Existence of the KY Transform). *If a function f belongs to class C^n at point x_0 and the connection function $g(t; \boldsymbol{\theta})$ has n independent parameters, then $\nabla_{KY}[f](x_0)$ exists.*

Theorem 3.13 (Rate of Convergence). *Under appropriate regularization conditions:*

$$\|\boldsymbol{\theta}(h) - \nabla_{KY}[f](x_0)\| = O(h^\alpha) \tag{24}$$

where $\alpha > 0$ is the order of convergence, which depends on the connection function and the smoothness of function f.

3.6 Specific Calculation Example

Example 3.1 (KY-Polynomial Calculation (Improved)). *KY-Parabolic Transform for the function $f(x) = x^3$ at the point $x_0 = 1$:*
  Connection function: $g(t) = a_2(t-1)^2 + a_1(t-1) + a_0$
  Constraints:

$$\mathcal{C}_1(h): \quad a_0 = f(1) = 1 \tag{25}$$
$$\mathcal{C}_2(h): \quad a_2 h^2 + a_1 h + a_0 = f(1+h) = (1+h)^3 \tag{26}$$

Expanding $(1+h)^3 = 1 + 3h + 3h^2 + h^3$:

$$a_2 h^2 + a_1 h + 1 = 1 + 3h + 3h^2 + h^3 \tag{27}$$

By comparing coefficients:

$$a_1(h) = 3 + h^2 \tag{28}$$
$$a_2(h) = 3 + h \tag{29}$$

Applying the regularization condition $\mathcal{R}_P^2$ (in this case, the solution is naturally unique):

$$\nabla_P^2[x^3](1) = \lim_{h \to 0}(3 + h, 3 + h^2, 1) = (3, 3, 1) \tag{30}$$

Interpretation of results:

- *Curvature coefficient = 3 = $\frac{f''(1)}{2!} = \frac{6}{2}$*

- *Slope = 3 = $f'(1)$*

- *Function value = 1 = $f(1)$*

3.7 Chain Rule

One of the most important properties of the KY Transform is the chain rule for composite functions.

Theorem 3.14 (KY Chain Rule). *For a composite function $h(x) = f(g(x))$, under appropriate regularization conditions:*

$$\nabla_{KY}[f \circ g](x_0) = \nabla_{KY}[f](g(x_0)) \otimes \nabla_{KY}[g](x_0) \tag{31}$$

where $\otimes$ is the KY chain operator, defined according to the connection function type.

3.7.1 Linear-Linear Chain Rule

Theorem 3.15 (Linear-Linear Chain Rule). *If $\nabla_L[f](u_0) = (a_f, b_f)$ and $\nabla_L[g](x_0) = (c_g, d_g)$, then:*

$$\nabla_L[f \circ g](x_0) = (a_f \cdot c_g, a_f \cdot d_g + b_f) \tag{32}$$

This is consistent with the traditional chain rule $[f \circ g]'(x_0) = f'(g(x_0)) \cdot g'(x_0)$.

3.8 Relationship with Conventional Derivatives

The KY Transform has a clear relationship with traditional derivatives and higher-order derivatives:

Theorem 3.16 (Relationship with Derivatives). *When function f is sufficiently smooth at point x_0:*

$$f'(x_0) = [\nabla_L[f](x_0)]_1 = [\nabla_P^2[f](x_0)]_1 \tag{33}$$
$$f''(x_0) = 2 \cdot [\nabla_P^2[f](x_0)]_2 \tag{34}$$
$$[\nabla_P^n[f](x_0)]_k = \frac{f^{(k)}(x_0)}{k!} \quad \text{(Taylor coefficient)} \tag{35}$$

3.9 Computational Caveats and Numerical Implementation

Key principles for correctly calculating the KY Transform:

1. **Consistent Application of Regularization Conditions**: Use the same regularization condition for all calculations.

2. **Matching Base Points**: For composite functions, it must be that $u_0 = g(x_0)$.

3. **Standard Form of Connection Function**: Use an expansion centered at the base point, in the form $(t - x_0)$.

4. **Numerical Stability**: Minimize error by selecting an appropriate value for h.

5. **Convergence Verification**: Numerically confirm the existence of the limit.

Remark 3.1 (Considerations for Numerical Implementation). *In practice, the following adaptive strategies are used:*

- *Adaptive reduction of h: $h_k = h_0 \cdot \rho^k$ (where $0 < \rho < 1$).*

- *Convergence criterion: $\|\boldsymbol{\theta}(h_{k+1}) - \boldsymbol{\theta}(h_k)\| < \epsilon$.*

- *Application of optimization algorithms for the numerical implementation of regularization conditions.*

Through such mathematical rigor and computational considerations, the KY Transform can be utilized as a powerful and consistent analytical tool that encompasses traditional differentiation.

4 Methodology

4.1 Applying KY Transform to Token Sequences

We can apply the KY Transform by interpreting a token sequence as a discretely sampled function. If we interpret the sequence $[token_1, token_2, \ldots, token_n]$ as a function $f(x)$, then $f(1) = token_1, f(2) = token_2, \ldots$.

The window $[token_{i-k}, token_{i-k+1}, \ldots, token_{i+k}]$ can be seen as a local sampling of the function, and the connection function $g(t; \theta)$ becomes a learnable function connecting these points. The result of the KY Transform, θ, represents the "dynamic characteristics" of that interval.

By physical intuition, if we apply a parabolic connection function $g(t) = a(t - t_0)^2 + b(t - t_0) + c$ to a position function $s(t)$, the KY result (a, b, c) corresponds to $(\text{acceleration}/2, \text{velocity}, \text{position})$, respectively. Similarly, in a token sequence, we can extract semantic dynamics such as the rate of grammatical change, the strength of semantic connections, and the contextual baseline.

4.2 Neural KY-Attention

Neural KY-Attention extends traditional self-attention to dynamically model complex relationships between tokens through a learnable connection function. Each attention head has an independent connection function neural network, allowing it to learn different kinds of relationships.

4.2.1 Connection Network Structure

The connection function for each attention head is implemented as the following neural network:

$$g_h(t) = \text{MLP}_h(t) \tag{36}$$
$$= W_3^{(h)} \cdot \text{GELU}(W_2^{(h)} \cdot \text{GELU}(W_1^{(h)} \cdot t + b_1^{(h)}) + b_2^{(h)}) + b_3^{(h)} \tag{37}$$

where $t \in [0, 1]$ is the normalized positional information, and h is the head index.

4.2.2 Algorithm Steps

The computation process of Neural KY-Attention is as follows:

Step 1: Standard Attention Calculation For an input sequence $X \in \mathbb{R}^{L \times d}$:

$$Q = XW^Q, \quad K = XW^K, \quad V = XW^V \tag{38}$$
$$Q, K, V \in \mathbb{R}^{L \times d} \rightarrow \mathbb{R}^{B \times H \times L \times d_h} \tag{39}$$

Step 2: Local Window Generation For a window size W, add padding and create a sliding window:

$$K_{\text{pad}} = \text{Pad}(K, W/2) \tag{40}$$
$$K_{\text{win}} = \text{Unfold}(K_{\text{pad}}, W) \in \mathbb{R}^{B \times H \times L \times W \times d_h} \tag{41}$$

The same is applied to V.

Step 3: Base Attention Scores Attention scores with tokens within the window at each position i:

$$\text{scores}_{i,j} = \frac{Q_i \cdot K_{\text{win}_{i,j}}}{\sqrt{d_h}} \tag{42}$$

Softmax normalization:

$$\alpha_{i,j} = \frac{\exp(\text{scores}_{i,j})}{\sum_{k=1}^{W} \exp(\text{scores}_{i,k})} \tag{43}$$

Step 4: Applying the Connection Function Connection weights for relative positions within the window:

$$t_j = \frac{j-1}{W-1}, \quad j = 1, 2, \ldots, W \tag{44}$$

$$c_{h,j} = g_h(t_j) = \text{MLP}_h(t_j) \tag{45}$$

$$\beta_{h,j} = \frac{\exp(c_{h,j})}{\sum_{k=1}^{W} \exp(c_{h,k})} \tag{46}$$

Step 5: Final Weight Combination Combination of attention scores and connection weights:

$$w_{i,j}^{(h)} = \alpha_{i,j} \cdot \beta_{h,j} \tag{47}$$

Normalization:

$$\tilde{w}_{i,j}^{(h)} = \frac{w_{i,j}^{(h)}}{\sum_{k=1}^{W} w_{i,k}^{(h)} + \epsilon} \tag{48}$$

Step 6: Output Calculation Final output:

$$\text{Output}_i^{(h)} = \sum_{j=1}^{W} \tilde{w}_{i,j}^{(h)} \cdot V_{\text{win}_{i,j}} \tag{49}$$

5 Theoretical Analysis

5.1 Expressive Power Analysis

Theorem 5.1 (Inclusion of Standard Attention). *Neural KY-Attention includes Standard Self-Attention as a special case.*

Proof. If we set the connection function $g_h(t) = c$ (a constant function):

1. Connection weights: $\beta_{h,j} = \frac{\exp(c)}{\sum_{k=1}^{W} \exp(c)} = \frac{1}{W}$ (uniform)

2. Final weights: $w_{i,j}^{(h)} = \alpha_{i,j} \cdot \frac{1}{W}$

3. Output: $\text{Output}_i^{(h)} = \frac{1}{W} \sum_{j=1}^{W} \alpha_{i,j} \cdot V_{\text{win}_{i,j}}$

In the limit where $W \to L$ (the entire sequence) and without padding:

$$\text{Output}_i^{(h)} = \sum_{j=1}^{L} \alpha_{i,j} \cdot V_j$$

This is identical to standard self-attention. $\qquad \square$

Theorem 5.2 (Expanded Expressive Power). *Neural KY-Attention can express richer relationships than Standard Attention.*

Proof. For the class of connection functions $\mathcal{G} = \{g : [0, 1] \to \mathbb{R} | g$ is a learnable function$\}$, the uniform weighting of Standard Attention is a special case where $g(t) = c$.

For any non-constant connection function $g^* \in \mathcal{G}$, we can generate position-specific weights as follows:

$$\beta^*_{h,j} = \frac{\exp(g^*(t_j))}{\sum_{k=1}^{W} \exp(g^*(t_k))} \neq \frac{1}{W}$$

This provides adaptive weighting based on position, modeling complex positional dependencies that Standard Attention cannot express. □

Theorem 5.3 (Universal Approximation). *Given a sufficient number of heads and connection function capacity, Neural KY-Attention can approximate any attention function.*

Proof. Let's define the function space as:

$$\mathcal{F} = \{f : \mathbb{R}^{L \times d} \to \mathbb{R}^{L \times d} | f \text{ is an attention transformation}\}$$

By the Universal Approximation Theorem [6], a sufficiently wide MLP can approximate any continuous function to an arbitrary precision.

For any target attention function $f^* \in \mathcal{F}$ and a given $\epsilon > 0$:

Step 1: Implement the connection function g_h of each head h as an MLP with sufficient capacity.

Step 2: There exist a sufficient number of heads H and appropriate combination weights λ_h such that:

$$\sup_{X \in \mathcal{X}} \left\| \sum_{h=1}^{H} \lambda_h \cdot \text{KY-Attn}_h(X) - f^*(X) \right\|_F < \epsilon$$

Step 3: If we bound the MLP approximation error of each g_h by δ, the total approximation error is proportional to $H \cdot \delta$. Thus, by selecting an appropriate H and MLP capacity, any ϵ can be achieved. □

5.2 Complexity Analysis

Theorem 5.4 (Time Complexity). *The time complexity of Neural KY-Attention is $O(LWH)$.*

Proof. Analyzing the complexity of each step:
Preprocessing Stage:

- Query, Key, Value calculation: $O(Ld^2 H)$

- Window generation (Padding + Unfold): $O(LW \cdot d)$

Core Operations:

- Attention score calculation: Dot product with W tokens in the window for each position i.
$$O(L \cdot W \cdot d_h \cdot H) = O(LWd) \text{ (since } d_h = d/H)$$

- Connection function calculation: MLP operation for each head.

$$O(H \cdot W \cdot C) = O(HW) \text{ (where } C \text{ is MLP operation constant)}$$

- Final weight calculation and output: $O(LWH)$

Excluding constant terms, the dominant term is $O(LWH)$. Compared to the $O(L^2H)$ of Standard Attention, this is a significant improvement when $W \ll L$. $\square$

Theorem 5.5 (Space Complexity). *The space complexity of Neural KY-Attention is $O(LWH)$.*

Proof. Analyzing the main memory usage:

$$\text{Input storage} : O(Ld) \tag{50}$$
$$\text{Window tensor} : O(L \cdot W \cdot H \cdot d_h) = O(LWd) \tag{51}$$
$$\text{Attention scores} : O(LWH) \tag{52}$$
$$\text{Connection weights} : O(WH) \tag{53}$$
$$\text{Intermediate results} : O(LWH) \tag{54}$$
$$\text{Output} : O(Ld) \tag{55}$$

The total space complexity is $O(LWd + LWH) = O(LWH)$ (since typically $H \geq d/H$).

Compared to storing the full attention matrix $O(L^2H)$ in Standard Attention, this significantly improves memory efficiency when $W \ll L$. $\square$

5.3 Convergence Analysis

Theorem 5.6 (Stability of Connection Function Learning). *Under appropriate initialization and regularization conditions, the learning of the connection function converges stably.*

Proof. Let's assume the connection function $g_h : [0, 1] \to \mathbb{R}$ is L-Lipschitz continuous:

$$|g_h(t_1) - g_h(t_2)| \leq L|t_1 - t_2|, \quad \forall t_1, t_2 \in [0, 1]$$

Gradient Bound: The gradient of the loss function $\mathcal{L}$ with respect to the connection function is:

$$\left\| \frac{\partial \mathcal{L}}{\partial g_h} \right\| \leq C \cdot L \cdot \|X\|_F \cdot \|\nabla_{\text{output}} \mathcal{L}\|$$

where C is the Lipschitz constant of the softmax function.
Stability Conditions:

1. **Initialization**: Xavier/He [4, 5] initialization ensures $\mathbb{E}[\|g_h(0)\|^2] = O(1)$.

2. **Normalization**: Applying Layer Normalization [1] reduces internal covariate shift.

3. **Gradient Clipping**: $\|\nabla g_h\| \leq \tau$ (for a threshold τ).

4. **Learning Rate Scheduling**: Adaptive learning rate adjustment.

Under these conditions, the SGD update satisfies:

$$\mathbb{E}[\|\theta_{t+1} - \theta^*\|^2] \leq (1 - \mu\eta)\mathbb{E}[\|\theta_t - \theta^*\|^2] + \sigma^2\eta^2$$

where μ is the strong convexity constant and σ^2 is the variance of the gradient noise. Convergence is guaranteed by the Robbins-Siegmund theorem [12]. $\square$

Theorem 5.7 (Gradient Propagation Stability). *Neural KY-Attention mitigates the vanishing/exploding gradient problem compared to the conventional Transformer.*

Proof. **Locality Property**: Due to the window-based operation:

$$\frac{\partial \text{Output}_j}{\partial X_i} = \begin{cases} \text{non-zero} & \text{if } |i - j| \leq W/2 \\ 0 & \text{otherwise} \end{cases}$$

Gradient Bound:

$$\left\|\frac{\partial \mathcal{L}}{\partial X_i}\right\| \leq \sum_{|j-i|\leq W/2} \left\|\frac{\partial \mathcal{L}}{\partial \text{Output}_j}\right\| \cdot \left\|\frac{\partial \text{Output}_j}{\partial X_i}\right\|$$

Due to the softmax normalization of the connection function $\sum_{j=1}^{W} \beta_{h,j} = 1$:

$$\left\|\frac{\partial \text{Output}_j}{\partial X_i}\right\| \leq M_{\text{softmax}} \cdot M_{\text{connection}}$$

where $M_{\text{softmax}} \leq 1$, and $M_{\text{connection}}$ is the maximum slope of the connection function. Therefore, when $W \ll L$ and the connection function is appropriately regularized:

$$\left\|\frac{\partial \mathcal{L}}{\partial X_i}\right\| \leq O(W) \cdot M_{\text{connection}} \cdot \left\|\frac{\partial \mathcal{L}}{\partial \text{Output}}\right\|$$

This is much more stable than the $O(L)$ dependency in Standard Attention. $\square$

5.4 Multi-Scale Representation Learning

Each head in Neural KY-Attention learns a different connection pattern to achieve a multi-layered understanding of the sequence:

$$\text{Head } h_1: \quad \nabla^{\text{KY}}[\text{sequence}; g_1] \rightarrow \text{Short-range grammatical patterns} \qquad (56)$$

$$\text{Head } h_2: \quad \nabla^{\text{KY}}[\text{sequence}; g_2] \rightarrow \text{Mid-range semantic associations} \qquad (57)$$

$$\text{Head } h_3: \quad \nabla^{\text{KY}}[\text{sequence}; g_3] \rightarrow \text{Long-range discourse structure} \qquad (58)$$

$$\text{Head } h_4: \quad \nabla^{\text{KY}}[\text{sequence}; g_4] \rightarrow \text{Thematic consistency} \qquad (59)$$

Theorem 5.8 (Multi-Pattern Decomposability). *A Neural KY-Attention with H heads can decompose an input sequence into H different perspectives.*

Proof. Let's assume the connection function g_h of each head h is learned from a different function class:

$$g_h \in \mathcal{G}_h, \quad \mathcal{G}_i \cap \mathcal{G}_j = \emptyset \text{ for } i \neq j$$

The overall representation for a sequence S is:

$$\text{Repr}(S) = \bigoplus_{h=1}^{H} \text{KY-Transform}(S; g_h)$$

where $\bigoplus$ is the concatenation or combination operation of the head-wise representations.

Since each g_h is learned independently, the following holds:

$$\text{Mutual Information}(\text{KY-Transform}(S; g_i), \text{KY-Transform}(S; g_j)) \rightarrow \min$$

This implies that each head extracts different, mutually complementary patterns. $\quad\square$

6 Experiments

6.1 Experimental Setup

We evaluated the performance of the Neural KY-Attention architecture on the language modeling task using the WikiText-2 dataset [8]. The hyperparameters used in the experiment are as follows:

- Model dimension: $d_{\text{model}} = 256$

- Number of layers: $L = 4$

- Number of attention heads: $H = 8$

- Feed-forward dimension: $d_{\text{ff}} = 1024$

- Maximum sequence length: 256

- Window size: $W = 15$ (applied only to Neural KY-Attention)

- Batch size: 16

- Training epochs: 3

- Learning rate: 5×10^{-4} (AdamW optimizer [7])

- Weight decay: 0.01

- Dropout: 0.1

The experiments were conducted in a CUDA-supported GPU environment, and the text was tokenized using the GPT-2 tokenizer [10]. A total of 5,000 training samples and 1,000 validation samples were used for the experiment.

Table 1: Performance Comparison between Neural KY-Attention and Standard Transformer

Metric	Neural KY Transformer	Standard Transformer
Final Validation Loss	6.3321	6.5739
Final Perplexity	562.31	716.14
Peak Training Memory (MB)	7,210.23	5,999.66
Inference Speed (samples/sec)	260.82	352.02
Inference Memory (MB)	3,122.14	3,118.38

6.2 Performance Comparison

We evaluated the performance of Neural KY-Attention through a comparative experiment with a standard Transformer model. The final experimental results are summarized in Table 1.

Neural KY-Attention showed the following performance improvements over the standard Transformer:

- **Perplexity**: A significant improvement of about 21.5%, with 562.31 vs 716.14.

- **Validation Loss**: An improvement of about 3.7%, with 6.33 vs 6.57.

- **Convergence Performance**: As shown in Figure 1, Neural KY-Attention shows faster and more stable convergence.

- **Memory Efficiency**: While inference memory usage is similar (3,122.14 MB vs 3,118.38 MB), it uses more memory during training (7,210.23 MB vs 5,999.66 MB).

- **Inference Speed**: The standard Transformer shows about 35.0% faster inference speed (352.02 vs 260.82 samples/sec).

6.3 Connection Pattern Analysis

The results of analyzing the connection function patterns learned by each attention head are shown in Figure 2. Analyzing the connection functions learned by the 8 heads of the first layer:

- **Head 1**: A concave pattern with gradually increasing weights toward later positions, emphasizing long-range dependencies.

- **Head 2**: A strong linear increase pattern, focusing on progressively distant positions.

- **Head 3**: A linear increase from negative to positive values, showing balanced attention distribution.

- **Head 4**: A steep linear increase pattern, strongly emphasizing later positions in the sequence.

- **Head 5**: A moderate linear increase, showing preference for distant dependencies.

- **Head 6**: A gradual increase with slight curvature, balancing local and distant attention.

- **Head 7**: A linear increase with consistent slope, uniformly distributing attention across positions.

- **Head 8**: A steep increase from negative to positive, showing strong positional bias.

These diverse connection patterns demonstrate that each head specializes in learning dependencies of different distances and positions. In particular, the learnable connection function adaptively discovers optimal position-wise weights according to the data characteristics, allowing the model to capture more nuanced positional relationships than fixed positional encodings.

7 Conclusion and Future Work

In this study, we proposed Neural KY-Attention, a novel attention mechanism based on the KY Transform theory. This approach overcomes the limitations of traditional pairwise interactions and dynamically models complex relationships between tokens through learnable connection functions.

The main achievements are as follows:

- Successful application of the KY Transform theory to the attention mechanism.

- Improved computational efficiency by achieving linear complexity $O(L \times W)$.

- Superior performance and faster convergence compared to the standard Transformer.

- Confirmation of specialized connection pattern learning in multi-head attention.

- Theoretical guarantees of expressiveness and convergence.

Future research directions include:

- Exploring various connection function architectures (CNN, RNN, Attention-based).

- An adaptive window size adjustment mechanism.

- Extension and application to other domains (vision, speech).

- Scalability research for longer sequences.

- Multi-scale modeling through hierarchical KY Transforms.

- Methods to improve the interpretability of connection functions.

In particular, the mathematical foundation of the KY Transform is applicable to other neural network components, holding the potential to create new variations of recurrent neural networks, convolutional neural networks, and more. It is also significant in terms of interpretability, as the learned patterns of the connection function can help us better understand the model's decision-making process.

A Appendix

A.1 Application Possibilities

A.1.1 Application to Computer Vision

Applying the KY Transform to sequences of image patches:

- **Spatial Connection Function**: Modeling relationships between patches in 2D space.

- **Scale Connection Function**: Learning relationships between different resolutions.

- **Channel Connection Function**: Modeling interactions between color/feature channels.

A.1.2 Application to Speech Processing

KY Transform on audio sequences:

- **Temporal Connection Function**: Modeling the temporal dynamics of speech.

- **Frequency Connection Function**: Relationships between frequencies in a spectrogram.

- **Prosodic Connection Function**: Learning patterns of intonation, stress, and rhythm.

A.1.3 Application to Graph Neural Networks

KY Transform on graph-structured data:

- **Path Connection Function**: Modeling the functional relationship of paths between nodes.

- **Neighborhood Connection Function**: Learning the dynamics of local neighborhood structures.

- **Hierarchical Connection Function**: Capturing multi-scale graph patterns.

B Detailed Comparison with Related Work

B.1 Comparison with Existing Efficient Attention Methods

B.2 Relationship with Positional Encoding Methods

The connection function of Neural KY-Attention has the following relationship with existing positional encoding methods:

- **Absolute PE**: Fixed positional information vs. learnable connection relationships.

- **Relative PE**: Based on relative distance vs. based on functional relationships.

Table 2: Comparison of Attention Methodologies

Method	Complexity	Expressiveness	Interpretability	Adaptability
Standard Attention	$O(L^2)$	High	Medium	Low
Linear Attention	$O(L)$	Low	Low	Low
Sparse Attention	$O(L\sqrt{L})$	Medium	Low	Low
Local Attention	$O(LW)$	Medium	Medium	Low
Neural KY-Attention	$O(LWH)$	High	High	High

- **RoPE**: Rotational transformation vs. general connection transformation.

- **ALiBi**: Linear decay vs. learnable arbitrary function.

Neural KY-Attention provides a generalized framework that can include all of these as special cases.

C   Limitations and Future Research Directions

C.1   Current Limitations

- **Hyperparameter Sensitivity**: The choice of window size and connection function structure has a large impact on performance.

- **Interpretation Complexity**: Interpreting the meaning of the learned connection functions is still challenging.

- **Very Long Sequences**: Currently optimized for medium-length sequences.

- **Domain Specificity**: A need to design connection functions suitable for each domain.

C.2   Future Research Directions

C.2.1   Automated Architecture Search

Applying Neural Architecture Search (NAS) [17, 11] to automatically search for the optimal structure of the connection function:

- Evolutionary algorithm-based search for connection function structure.

- Reinforcement learning for adaptive window size adjustment.

- Automatic discovery of optimal connection patterns for specific tasks.

C.2.2   Theoretical Extensions

- Information-theoretic analysis of the KY Transform.

- PAC-Bayes analysis of the expressive power and generalization ability of connection functions.

- Approximation theory research on various classes of connection functions.

C.2.3 Practical Applications

- Application to large-scale language models and study of scaling laws.

- Extension of cross-attention in multimodal models.

- Hardware optimization for real-time inference.

This research presents a new paradigm for the attention mechanism through the KY Transform theory, which is expected to provide an important theoretical basis for the future development of neural network architectures. In particular, the concept of a learnable connection function is a universal idea with great potential, applicable not only to attention but also to various other neural network components.

D   Figures

Figure 1: Training and validation loss curves for Neural KY Transformer vs. Standard Transformer.


Figure 2: Learned connection function patterns for the 8 attention heads in the first layer.

References

[1] Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E Hinton. Layer normalization. *arXiv preprint arXiv:1607.06450*, 2016.

[2] Iz Beltagy, Matthew E Peters, and Arman Cohan. Longformer: The long-document transformer. *arXiv preprint arXiv:2004.05150*, 2020.

[3] Krzysztof Choromanski, Valerii Likhosherstov, David Dohan, Xingyou Song, Andreea Gane, Tamas Sarlos, Peter Hawkins, Jared Davis, Afroz Mohiuddin, Lukasz Kaiser, et al. Rethinking attention with performers. *International Conference on Learning Representations*, 2021.

[4] Xavier Glorot and Yoshua Bengio. Understanding the difficulty of training deep feedforward neural networks. *Proceedings of the thirteenth international conference on artificial intelligence and statistics*, pages 249–256, 2010.

[5] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Delving deep into rectifiers: Surpassing human-level performance on imagenet classification. *Proceedings of the IEEE international conference on computer vision*, pages 1026–1034, 2015.

[6] Kurt Hornik, Maxwell Stinchcombe, and Halbert White. Multilayer feedforward networks are universal approximators. *Neural networks*, 2(5):359–366, 1989.

[7] Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *International Conference on Learning Representations*, 2015.

[8] Stephen Merity, Caiming Xiong, James Bradbury, and Richard Socher. Pointer sentinel mixture models. *International Conference on Learning Representations*, 2017.

[9] Ofir Press, Noah A Smith, and Mike Lewis. Train short, test long: Attention with linear biases enables input length extrapolation. *International Conference on Learning Representations*, 2021.

[10] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners, 2019.

[11] Esteban Real, Alok Aggarwal, Yanping Huang, and Quoc V Le. Regularized evolution for image classifier architecture search. *Proceedings of the AAAI conference on artificial intelligence*, 33(01):4780–4789, 2019.

[12] Herbert Robbins and Sutton Monro. *A stochastic approximation method*. 1951.

[13] Jianlin Su, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yunfeng Liu. Roformer: Enhanced transformer with rotary position embedding. *arXiv preprint arXiv:2104.09864*, 2021.

[14] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30:5998–6008, 2017.

[15] Sinong Wang, Belinda Z Li, Madian Khabsa, Han Fang, and Hao Ma. Linformer: Self-attention with linear complexity. *arXiv preprint arXiv:2006.04768*, 2020.

[16] Manzil Zaheer, Guru Guruganesh, Kumar Avinava Dubey, Joshua Ainslie, Chris Alberti, Santiago Ontanon, Philip Pham, Anirudh Ravula, Qifan Wang, Li Yang, et al. Big bird: Transformers for longer sequences. *Advances in neural information processing systems*, 33:17283–17297, 2020.

[17] Barret Zoph and Quoc V Le. Neural architecture search with reinforcement learning. *International Conference on Learning Representations*, 2017.